

Perspectives on Probabilistic Graphical Models

Dong Liu

*Information Science and Engineering
Department of Intelligent Systems
KTH - Royal Institute of Technology*





Profile page: <https://firsthandscientist.github.io/>



Slide is available at: <https://github.com/FirstHandScientist/phdthesis>

Why are probabilistic graphical models interesting?

RICH REPRESENTATIONS

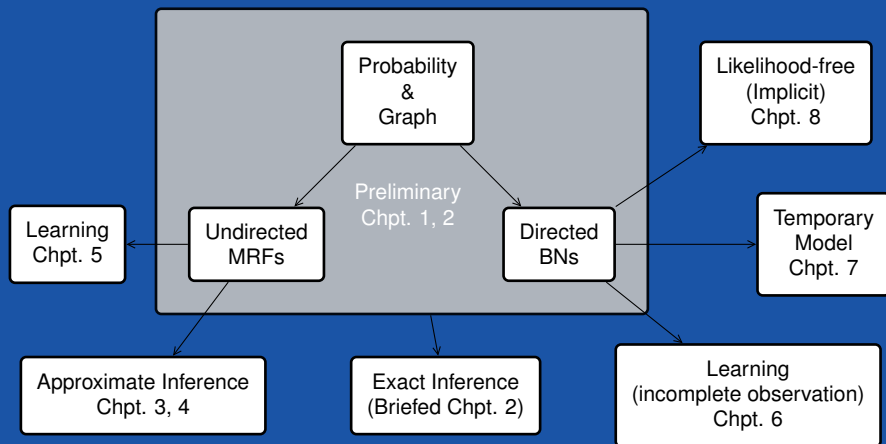


- Computer perception
- Error-control codes



- Computational biology
- Natural language processing
- etc.

A GUIDE TO THIS DISSERTATION



WHAT ARE PROBABILISTIC GRAPHICAL MODELS

Informally...

A PGM is a structured graph representation to encode

- Attributes of our interests in a system → variable nodes
- Relationship of these factors → structures of a graph


Intrinsic property: **reasoning with uncertainty**

WHAT ARE PROBABILISTIC GRAPHICAL MODELS

EXEMPLIFIED DEFINITIONS

A directed/undirected graph encoding dependencies/independencies of distribution $p(\mathbf{x}; \theta)$:


- A Generative model/BN is a directed graph (DAG)



$$p(\mathbf{x}; \theta) = \prod_{n=1}^N p(x_n | \underbrace{\mathcal{P}(x_n)}_{\text{parent nodes of } x_n})$$

the local functions are proper distributions

- An MRF denoted by an undirected graph $\mathcal{G}(\mathcal{V}, \mathcal{E})$




- The probability distribution (Gibbs distribution) is $p(\mathbf{x}; \theta) = \frac{1}{Z(\theta)} \prod_{a \in \mathcal{I}} \underbrace{\psi_a(\mathbf{x}_a; \theta_a)}_{\text{potential function is not necessarily a proper distribution}}$
- a indexes potential functions
 $\mathcal{I} = \{\psi_A, \psi_B, \dots, \psi_M\}$
- $Z(\theta) = \sum_{\mathbf{x}} \prod_a \psi_a(\mathbf{x}_a; \theta_a)$.

WHAT TO DO WITH GRAPHICAL MODELS

- Inference
 - Computing the likelihood of observed data.
 - Computing the marginals distribution $p(\mathbf{x}_A)$ over particular subset $A \subset \mathcal{V}$ of nodes
 - Computing the conditional distribution $p(\mathbf{x}_A | \mathbf{x}_B)$,
 - Computing the partition function or the Helmholtz free energy (for MRFs)
- Learning
 - To model or determine $p(\mathbf{x}; \theta)$.

Two key components interacting with each other:




WHAT IS THE STATE OF x ?

A TOY EXAMPLE

Assume that we are interested into the state of node i in an MRF, it can be answered by

- an empirical version, a collection of samples $\{x_i^n\}_{n=1}^N$, (sampling techniques)
- the probability $p(x_i)$




What is the state of x_4

WHAT IS THE STATE OF x ?

A TOY EXAMPLE

Assume that we are interested into the state of node i in an MRF, it can be answered by

- an empirical version, a collection of samples $\{x_i^n\}_{n=1}^N$, (sampling techniques)
- the probability $p(x_i)$



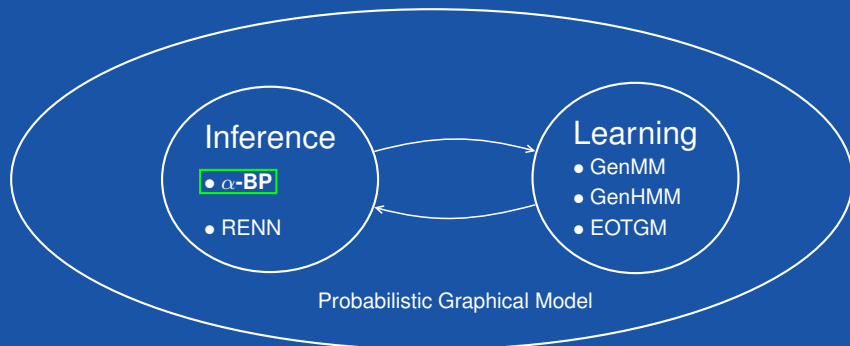
Mean Field and BP: *message in form of sample values* \rightarrow *message in form of belief*

- Propagating beliefs iteratively
- Queries by collected beliefs $\{m_i\}$.

Intuition from *Gibbs (variational) free energy*

$$F_V(b) = \text{KL}(b(\mathbf{x}) || p(\mathbf{x}; \theta)) - \log Z(\theta)$$


with trial $b(\mathbf{x})$. Instance: Bethe free energy.



ALTERNATIVE VIEW OF BP: α -BP


Ingredients:

- A pairwise Markov random field:
 $p(\mathbf{x}) \propto \prod_{s \in \mathcal{V}} \varphi_s(x_s) \prod_{(s,t) \in \mathcal{E}} \varphi_{st}(x_s, x_t)$
- A trial distribution: $q(\mathbf{x}) \propto \prod_{s \in \mathcal{V}} \tilde{\varphi}_s(x_s) \prod_{(s,t) \in \mathcal{E}} \tilde{\varphi}_{st}(x_s, x_t)$
 with factorization
 $\tilde{\varphi}_{s,t}(x_s, x_t) := m_{st}(x_t) m_{ts}(x_s)$
- A metric: α -Divergence

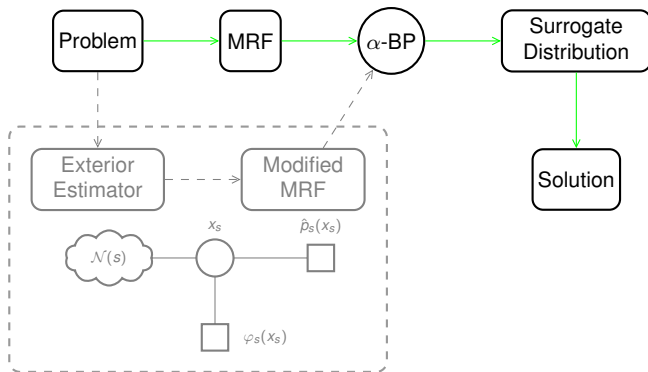


A factor graph representation
 $\mathcal{G}_F := (\mathcal{V} \cup \mathcal{F}, \mathcal{E}_F)$


Updating message via α -BP:



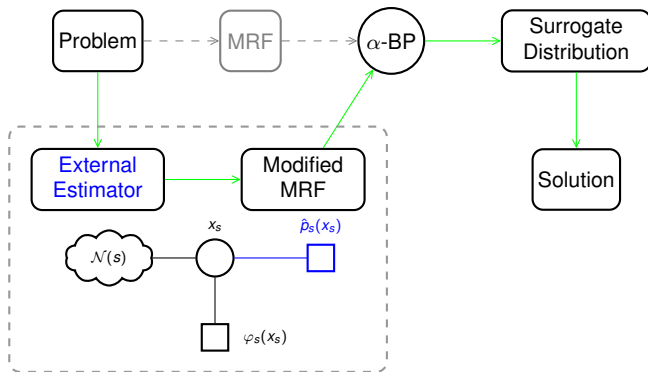
$$\underbrace{m_{ts}^{\text{new}}(x_s)}_{\text{new msg. to s}} \propto \underbrace{m_{ts}(x_s)^{1-\alpha_{ts}}}_{\text{old msg. to s}} \left[\sum_{x_t} \varphi_{ts}(x_t, x_s)^{\alpha_{ts}} \underbrace{m_{st}(x_t)^{1-\alpha_{ts}}}_{\text{old msg. to t}} \underbrace{\varphi_t(x_t) \prod_{w \in \mathcal{N}(t) \setminus s} m_{wt}(x_t)}_{\text{msg. from variable node } t \text{ to factor } \varphi_{st}} \right].$$



Updating message via α -BP:



$$\underbrace{m_{ts}^{\text{new}}(x_s)}_{\text{new msg. to s}} \propto \underbrace{m_{ts}(x_s)^{1-\alpha_{ts}}}_{\text{old msg. to s}} \left[\sum_{x_t} \varphi_{ts}(x_t, x_s)^{\alpha_{ts}} \underbrace{m_{st}(x_t)^{1-\alpha_{ts}}}_{\text{old msg. to t}} \underbrace{\varphi_t(x_t) \prod_{w \in \mathcal{N}(t) \setminus s} m_{wt}(x_t)}_{\text{msg. from variable node } t \text{ to factor } \varphi_{st}} \right].$$



INSIGHTS OF α -BP

Connection to standard BP

- $\alpha \rightarrow 1$
- α -divergence reduces to KL-divergence
- Update rule of α -BP reduces to

$$m_{ts}^{\text{new}}(x_s) \propto \sum_{x_t} \varphi_{st}(x_s, x_t) \varphi_t(x_t) \prod_{w \in \mathcal{N}(t) \setminus s} m_{wt}(x_t),$$
 which is standard BP update rule

Convergence

For an arbitrary pairwise Markov random field over binary variables, if the largest singular value of matrix $\mathbf{M}(\alpha, \theta)$ is less than one, α -BP converges to a fixed point. The associated fixed point is unique.

See Corollary 3.1 for relaxed condition where singular value computation is avoided.


What does that mean

- You can use α -BP as an alternative to (loopy) BP
- You can use matrix \mathbf{M} to check if you are guaranteed to get stable solution from α -BP

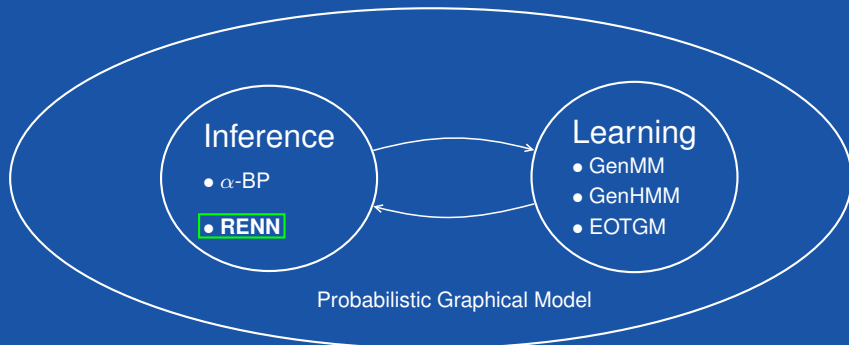
Matrix $\mathbf{M}(\alpha, \theta)$, size $|\mathcal{E}| \times |\mathcal{E}|$

Each element is either 0 or a function of α and potentials factors

SOME NUMERICAL RESULTS: APPLICATION CASE



Numerical results of α -BP: symbol error of MIMO detection.



CONTINUING: WHAT IS THE STATE OF x ?

YEDIDIA, FREEMAN, WEISS: A STEP TO GENERALIZATION

Message among variables & factors → message among regions

Generalized belief propagation (GBP) generalizes loopy BP

- usual better approximation than LBP
- higher complexity
- sensitive to scheduling of region messages

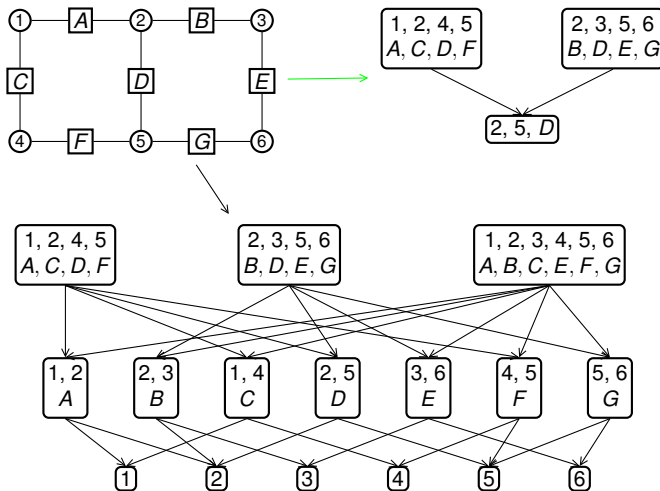
Approximating **variational free energy** $F_V(b)$ with trial b including $\{b_R\}$.

A *region* R is a set V_R of variables nodes and a set A_R of factor nodes, such that if a factor node ' a ' belongs to A_R , all the variables nodes neighboring a are in V_R .

A TOY EXAMPLE OF REGION GRAPHS

Factor graph representation of MRF (2-by-3 grid) with factor nodes.

MRF \rightarrow region graph:



- Clustering nodes
- level/layer-wise
- Hierarchical
- Msg. Scheduling
- ...
- See Section 4.1

RENN: REGION-BASED ENERGY NEURAL NETWORK

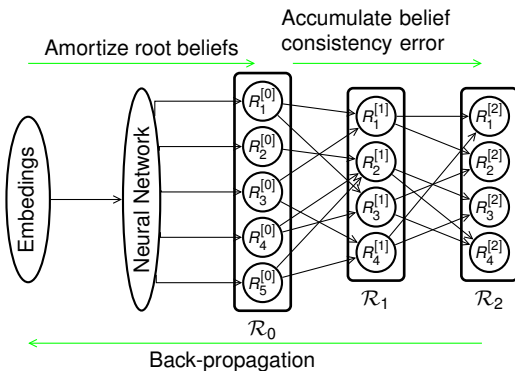

The **region-based free energy** of a region graph is

$$F_R(\mathcal{B}; \theta) = \sum_{R \in \mathcal{R}} \underbrace{c_R}_{\text{Counting number}} \underbrace{(\text{region average energy} - \text{region entropy})}_{\text{region free energy}},$$

RENN: REGION-BASED ENERGY NEURAL NETWORK


The **region-based free energy** of a region graph is

$$F_R(\mathcal{B}; \theta) = \sum_{R \in \mathcal{R}} \underbrace{C_R}_{\text{Counting number}} \underbrace{(\text{region average energy} - \text{region entropy})}_{\text{region free energy}},$$



RENN

Non-root belief:




Objective of RENN:

$$\min_{\text{parameter of NN}} \underbrace{\text{region-based free energy}(F_R)}_{\text{Accumulated over all regions}} + \underbrace{\text{penalty on belief inconsistency}}_{\text{Recursively computed via levels of region graph}}$$

RENN


Non-root belief:



Objective of RENN:

$$\min_{\text{parameter of NN}} \underbrace{\text{region-based free energy}(F_R)}_{\text{Accumulated over all regions}} + \underbrace{\text{penalty on belief inconsistency}}_{\text{Recursively computed via levels of region graph}}$$

RENN Inference:



RENN

INSIGHTS OF RENN

Generalization

Bethe free energy can be recovered from region-based free energy:

- two-level region graph representation
- constraint that each region can contain at most one factor node

Section 4.2.1

Attributes of RENN

- RENN requires neither sampling technique nor training data (ground-truth marginal probabilities) in performing inference tasks; **on-the-fly inference**
- RENN does gradient descent w.r.t. its neural network parameter instead of iterative message-passing, and returns approximation of marginal probabilities and partition estimation **in one-shot**
- No message propagation, thus **no need of message scheduling**
- Competitive performance and efficiency

SOME NUMERICAL COMPARISONS

Ising model: $p(\mathbf{x}; \theta) = \frac{1}{Z(\theta)} \exp(\sum_{(i,j) \in \mathcal{E}_F} J_{ij} x_i x_j + \sum_{i \in \mathcal{V}} h_i x_i)$, $\mathbf{x} \in \{-1, 1\}^N$,

- J_{ij} is the pairwise log-potential between node i and j , $J_{ij} \sim \mathcal{N}(0, 1)$
- h_i is the node log-potential for node i , $h_i \sim \mathcal{N}(0, \gamma^2)$

Inference on grid graph ($\gamma = 0.1$).

Metric	n	Mean Field	Loopy BP	Damped BP	GBP	Inference Net	RENN
ℓ_1 error	25	0.271 ± 0.051	0.086 ± 0.078	0.084 ± 0.076	0.057 ± 0.024	0.111 ± 0.072	0.049 ± 0.078
	100	0.283 ± 0.024	0.085 ± 0.041	0.062 ± 0.024	0.064 ± 0.019	0.074 ± 0.034	0.025 ± 0.011
	225	0.284 ± 0.019	0.100 ± 0.025	0.076 ± 0.025	0.073 ± 0.013	0.073 ± 0.012	0.046 ± 0.011
	400	0.279 ± 0.014	0.110 ± 0.016	0.090 ± 0.016	0.079 ± 0.009	0.083 ± 0.009	0.061 ± 0.009
Corre- lation ρ	25	0.633 ± 0.197	0.903 ± 0.114	0.905 ± 0.113	0.923 ± 0.045	0.866 ± 0.117	0.951 ± 0.112
	100	0.582 ± 0.112	0.827 ± 0.134	0.902 ± 0.059	0.899 ± 0.043	0.903 ± 0.049	0.983 ± 0.012
	225	0.580 ± 0.080	0.801 ± 0.078	0.863 ± 0.088	0.869 ± 0.037	0.873 ± 0.037	0.949 ± 0.022
	400	0.596 ± 0.054	0.779 ± 0.059	0.822 ± 0.047	0.852 ± 0.024	0.841 ± 0.028	0.912 ± 0.025
log Z error	25	2.512 ± 1.060	0.549 ± 0.373	0.557 ± 0.369	0.169 ± 0.142	0.762 ± 0.439	0.240 ± 0.140
	100	13.09 ± 2.156	1.650 ± 1.414	1.457 ± 1.365	0.524 ± 0.313	2.836 ± 2.158	1.899 ± 0.495
	225	29.93 ± 4.679	3.348 ± 1.954	3.423 ± 2.157	1.008 ± 0.653	3.249 ± 2.058	4.344 ± 0.813
	400	51.81 ± 4.706	5.738 ± 2.107	5.873 ± 2.211	1.750 ± 0.869	3.953 ± 2.558	7.598 ± 1.146

- ℓ_1 error of beliefs v.s. true
- correlation ρ between true and approximate marginals,
- log Z error, true v.s. free energy approximation.

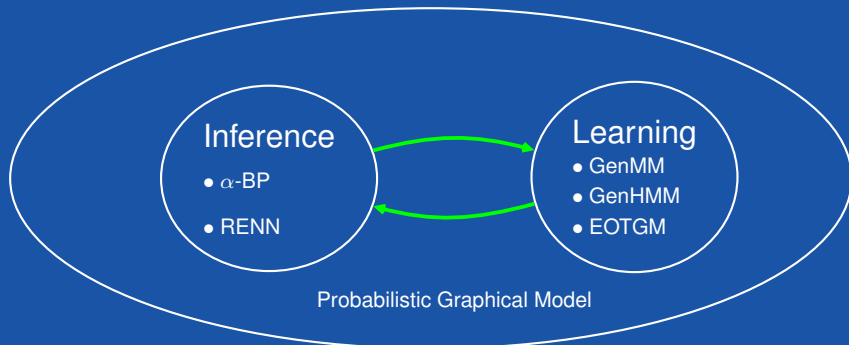
SOME NUMERICAL COMPARISONS

RICHER COMPARISONS

Inference on grid and complete graphs.

		Metric		Mean Field	Loopy BP	Damped BP	GBP	Inference Net	RENN
High Temp -erature	Complete graph N=16	ℓ_1 -	$\gamma = 1$	0.273 ± 0.086	0.239 ± 0.059	0.239 ± 0.059	0.260 ± 0.086	0.249 ± 0.067	0.181 ± 0.092
		error	$\gamma = 4$	0.197 ± 0.049	0.181 ± 0.035	0.180 ± 0.034	0.210 ± 0.070	0.174 ± 0.030	0.125 ± 0.050
	$J_{ij} \sim \mathcal{N}(0, 1)$ $h_i \sim \mathcal{N}(0, \gamma^2)$	log Z	$\gamma = 1$	20.66 ± 5.451	178.7 ± 22.18	178.9 ± 21.88	153.3 ± 25.29	213.6 ± 12.75	14.41 ± 4.135
		error	$\gamma = 4$	10.74 ± 7.385	565.7 ± 73.33	566.1 ± 73.13	106.0 ± 54.43	588.3 ± 62.58	14.72 ± 4.155
Low Temp -erature	Grid graph N=100	ℓ_1	5	0.257 ± 0.065	0.115 ± 0.071	0.120 ± 0.073	0.250 ± 0.024	0.164 ± 0.036	0.100 ± 0.046
		error	15	0.328 ± 0.068	0.228 ± 0.088	0.267 ± 0.147	0.303 ± 0.026	0.279 ± 0.024	0.207 ± 0.054
	$J_{ij} \sim \mathcal{U}(-u, u)$ $h_i \sim \mathcal{U}(-1, 1)$	log Z	5	42.65 ± 17.86	7.346 ± 7.744	5.444 ± 4.811	8.369 ± 7.401	65.60 ± 8.786	11.34 ± 4.724
		error	15	164.9 ± 56.07	58.40 ± 41.36	101.9 ± 54.31	23.10 ± 15.06	224.3 ± 25.52	78.85 ± 15.08

Low temperature setting translates to high variance of coupling strength between nodes (larger variance of J_{ij}).



INFERENCE ROUTINE IN LEARNING

What is θ in $p(\mathbf{x}; \theta) = \frac{1}{Z(\theta)} \prod_a \psi_a(\mathbf{x}_a; \theta_a)$?

A direct view:

$$\max_{\theta} \log p(\mathbf{x}; \theta) = \max_{\theta} \sum_a \log \psi_a(\mathbf{x}_a; \theta_a) \quad \underbrace{-\log Z(\theta)}_{\text{can be est. by min } F_V},$$

An alternative view:

$$\frac{\partial \log p(\mathbf{x}; \theta)}{\partial \theta_a} = \frac{\partial \log \varphi_a(\mathbf{x}_a; \theta_a)}{\partial \theta_a} - \underbrace{\mathbb{E}_{p(\mathbf{x}_a; \theta)} \left[\frac{\partial \log \varphi_a(\mathbf{x}_a; \theta_a)}{\partial \theta_a} \right]}_{\text{can be est. by beliefs}}.$$

Remark:

- This essentially requires estimation of partition function or marginal probabilities.
- Stationary points translate into moment matching.

INFERENCE ROUTINE IN LEARNING

What is θ in $p(\mathbf{x}; \theta) = \frac{1}{Z(\theta)} \prod_a \psi_a(\mathbf{x}_a; \theta_a)$?

A direct view:

$$\max_{\theta} \log p(\mathbf{x}; \theta) = \max_{\theta} \sum_a \log \psi_a(\mathbf{x}_a; \theta_a) - \underbrace{\log Z(\theta)}_{\text{can be est. by min } F_V},$$

An alternative view:

$$\frac{\partial \log p(\mathbf{x}; \theta)}{\partial \theta_a} = \frac{\partial \log \varphi_a(\mathbf{x}_a; \theta_a)}{\partial \theta_a} - \underbrace{\mathbb{E}_{p(\mathbf{x}_a; \theta)} \left[\frac{\partial \log \varphi_a(\mathbf{x}_a; \theta_a)}{\partial \theta_a} \right]}_{\text{can be est. by beliefs}}.$$

Remark:

- This essen
- Stationary



- Two modules are not necessarily coupled
- Each module may be replaced by another algorithm while the other one remains.

bilities.

LEARNING MRFs

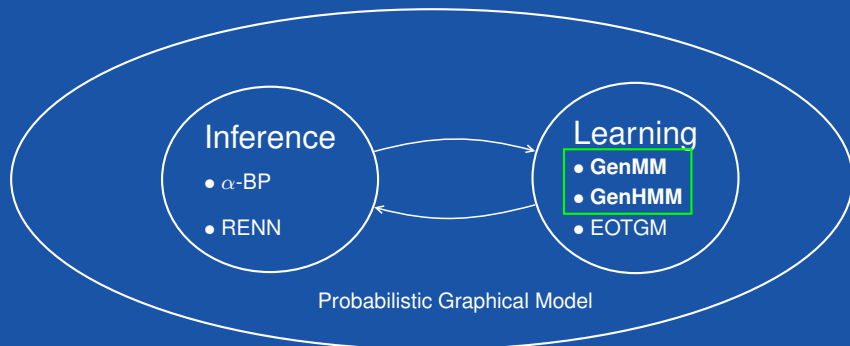
WHAT IS θ IN $p(\mathbf{x}; \theta)$?

Table of negative log-likelihood of learned MRFs

N	True	Exact	Mean Field	Loopy BP	Damped BP	GBP	Inference Net	RENN
Grid Graph								
25	9.000	9.004	9.811	9.139	9.196	10.56	9.252	9.048
100	19.34	19.38	23.48	19.92	20.02	28.61	20. 29	19.76
225	63.90	63.97	69.01	66.44	66.25	92.62	68.15	64.79
Complete Graph								
9	3.276	3.286	9.558	5.201	5.880	10.06	5.262	3.414
16	4.883	4.934	28.74	13.64	18.95	24.45	13.77	5.178

Average consumed time per epoch (unit: second) for two MRF learning cases.

	Mean Field	Loopy BP	Damped BP	GBP	Inference Net	RENN
Grid \mathcal{G} , $N=225$	40.09	335.1	525.1	12.37	19.49	16.03
Complete \mathcal{G} , $N=16$	2.499	12.40	5.431	1.387	0.882	2.262



INCOMPLETE OBSERVATION

Partial observation: $\mathbf{x} = [\underbrace{\mathbf{x}_U}_{\text{Unobserved}}, \underbrace{\mathbf{x}_O}_{\text{Observed}}]$

$$l(\mathbf{x}_O; \theta) = \log \sum_{\mathbf{x}_U} p(\mathbf{x}_U, \mathbf{x}_O; \theta) = \underbrace{\sum_{\mathbf{x}_U} \underbrace{\log Z(\mathbf{x}_O; \theta)}_{\substack{\log Z(\theta) \\ \text{0 in DAGs}}} \underbrace{\tilde{p}(\mathbf{x}; \theta)}_{\text{generalize}}}_{\text{both may be est. by free energy minimization in MRFs}} - \log Z(\theta)$$

Connect Free Energy to Evidence Lower Bounder:

$$\begin{aligned} l(\mathbf{x}_O; \theta) &\geq - \underbrace{F_V(q(\mathbf{x}_U | \mathbf{x}_O))}_{\text{Variational Free Energy}} - \log Z(\theta) \\ &= \mathbb{E}_{q(\mathbf{x}_U | \mathbf{x}_O)} \left[\log \frac{p(\mathbf{x}_U, \mathbf{x}_O; \theta)}{q(\mathbf{x}_U | \mathbf{x}_O)} \right] \\ &= \underbrace{\mathbb{E}_{q(\mathbf{x}_U | \mathbf{x}_O)} [\log p(\mathbf{x}_U, \mathbf{x}_O; \theta)] + H(q(\mathbf{x}_U | \mathbf{x}_O))}_{\text{Evidence Lower Bound } F(q, \theta)} \end{aligned}$$

Intuition of maximizing $F(q, \theta)$

- Maximizing (incomplete) likelihood
- Minimizing free energy

This gives raise of EM as a coordinate ascent method:

E step : $q^{(t+1)} = \underset{q}{\operatorname{argmax}} F(q, \theta^{(t)})$,

M step : $\theta^{(t+1)} = \underset{\theta}{\operatorname{argmax}} F(q^{(t+1)}, \theta)$.

GENERATOR MIXED MODEL

EQUIPPING EM WITH NORMALIZING FLOWS

- Ideal case: The underline true $p^*(\mathbf{x})$ is in hypothesis space \mathcal{H} , i.e. $p^*(\mathbf{x}) \in \mathcal{H}$.
- Out of reach: Test $p^*(\mathbf{x}) \stackrel{?}{\in} \mathcal{H}$
- Luckily, what is at our hands is:

\mathcal{H} is large \rightarrow candidate $p(\mathbf{x}; \theta)$ is flexible

This brings up the finite **mixture** models.

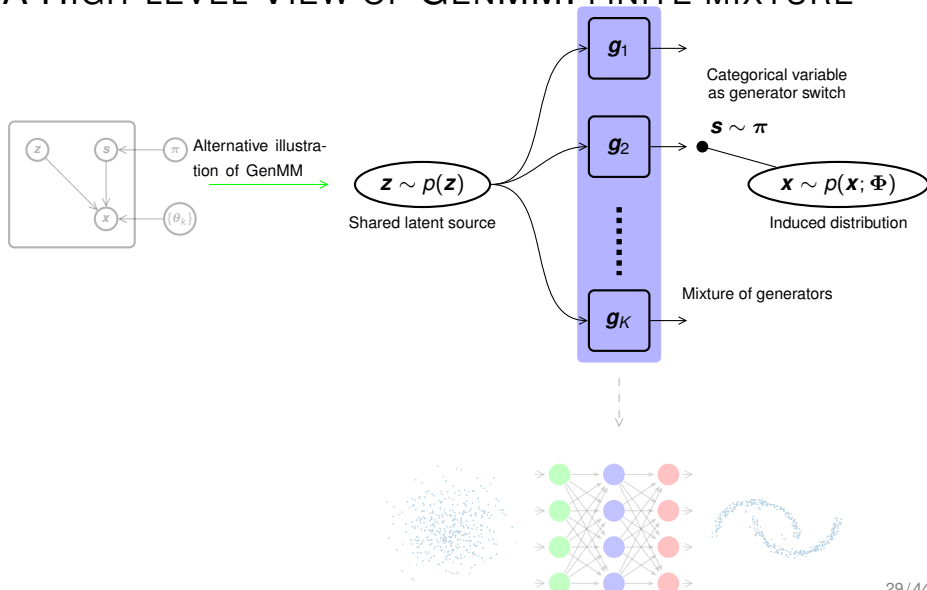
$$p(\mathbf{x}; \Theta) = \sum_{k=1}^K \pi_k p_k(\mathbf{x}) = \sum_{k=1}^K \pi_k p(\underbrace{g(\mathbf{z}; \theta_k)}_{\text{Variable change via generator } g})$$

Variable change
via generator g


What to expect from GenMM:

- Flexible and expressive model, enlarging hyperspace \mathcal{H}
- Tractable likelihood
- Compatible with typical statistical models
- Compatible with NN tools/frameworks
- Scale to high-dimensional structured data
- Efficient in sampling (data generation)
- ...

A HIGH-LEVEL VIEW OF GENMM: FINITE MIXTURE



A HIGH-LEVEL VIEW OF GENMM: FLOW GEARS




When the k -th generator is selected, i.e., $s_k = 1$ and $s_{k'} = 0$ for $k' \neq k$, say $\tilde{\mathbf{x}} = \mathbf{x}|_{s_k=1}$. By following the [change of variable rule](#)

$$\underbrace{p(\tilde{\mathbf{x}})|_{\tilde{\mathbf{x}}=\tilde{\mathbf{g}}(\mathbf{z})}}_{\text{Induced distribution}} = \underbrace{p(\mathbf{z})}_{\substack{\text{Assumed known distribution} \\ \text{easy to sample}}} \cdot \underbrace{\left| \det \left(\frac{\partial \mathbf{z}}{\partial \tilde{\mathbf{x}}} \right) \right|}_{\substack{\text{Computational load} \\ \text{depends on the mapping}}}.$$

A toy example:

$$\text{Gaussian linear transform: } Z \sim N(0, 1) \xrightarrow{X=\sigma \cdot Z + \mu} X \sim N(\mu, \sigma)$$

A HIGH-LEVEL VIEW OF GENMM: FLOW GEARS




When the k -th generator is selected, i.e., $s_k = 1$ and $s_{k'} = 0$ for $k' \neq k$, say $\tilde{\mathbf{x}} = \mathbf{x}|_{s_k=1}$. By following the [change of variable rule](#)

$$\underbrace{p(\tilde{\mathbf{x}})|_{\tilde{\mathbf{x}}=\tilde{\mathbf{g}}(\mathbf{z})}}_{\text{Induced distribution}} = \underbrace{p(\mathbf{z})}_{\substack{\text{Assumed known distribution} \\ \text{easy to sample}}} \cdot \underbrace{\left| \det \left(\frac{\partial \mathbf{z}}{\partial \tilde{\mathbf{x}}} \right) \right|}_{\substack{\text{Computational load} \\ \text{depends on the mapping}}} .$$

Powering it with a L -layer neural network implementation:

$$\mathbf{z} = \mathbf{h}_0 \xrightleftharpoons[\tilde{\mathbf{f}}_1]{\tilde{\mathbf{g}}_1} \mathbf{h}_1 \xrightleftharpoons[\tilde{\mathbf{f}}_2]{\tilde{\mathbf{g}}_2} \dots \xrightleftharpoons[\tilde{\mathbf{f}}_L]{\tilde{\mathbf{g}}_L} \mathbf{x} = \mathbf{h}_L$$

A HIGH-LEVEL VIEW OF GENMM: FLOW GEARS



When the k -th generator is selected, i.e., $s_k = 1$ and $s_{k'} = 0$ for $k' \neq k$, say $\tilde{\mathbf{x}} = \mathbf{x}|_{s_k=1}$. By following the [change of variable rule](#)

$$\underbrace{p(\tilde{\mathbf{x}})|_{\tilde{\mathbf{x}}=\tilde{\mathbf{g}}(\mathbf{z})}}_{\text{Induced distribution}} = \underbrace{p(\mathbf{z})}_{\substack{\text{Assumed known distribution} \\ \text{easy to sample}}} \cdot \underbrace{\left| \det \left(\frac{\partial \mathbf{z}}{\partial \tilde{\mathbf{x}}} \right) \right|}_{\substack{\text{Computational load} \\ \text{depends on the mapping}}}.$$

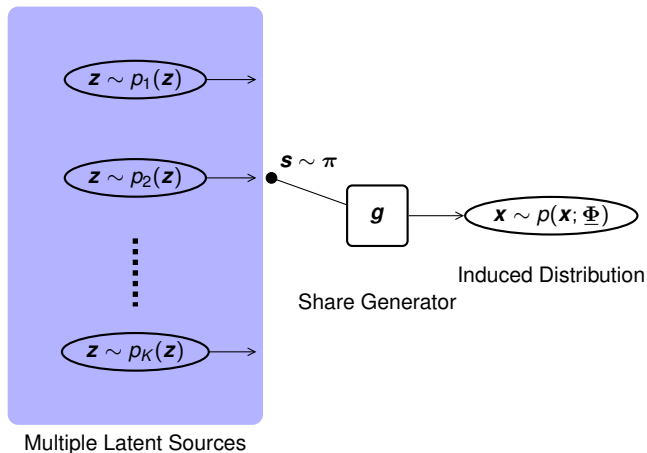
Layer structure:

- RealNVP
- Glow
- ODE

< - - - -

$$\mathbf{z} = \mathbf{h}_0 \xleftrightarrow[\tilde{\mathbf{f}}_1]{\tilde{\mathbf{g}}_1} \mathbf{h}_1 \xleftrightarrow[\tilde{\mathbf{f}}_2]{\tilde{\mathbf{g}}_2} \dots \xleftrightarrow[\tilde{\mathbf{f}}_L]{\tilde{\mathbf{g}}_L} \mathbf{x} = \mathbf{h}_L$$


LATMM: ALTERNATIVE MIXTURE



SAMPLING EXAMPLES



Generated samples by GenMM and LatMM.




Interpolation in latent space

APPLICATION TO CLASSIFICATION TASKS

Application to classification with maximum likelihood. Test Accuracy Table of GenMM for Classification Task

Dataset	K=1	K=2	K=3	K=4	K=10	K=20	State Of Art
Letter	0.9459	0.9513	0.9578	0.9581	0.9657	0.9674	0.9582
Satimage	0.8900	0.8975	0.9045	0.9085	0.9105	0.9160	0.9090
Norb	0.9184	0.9257	0.9406	0.9459	0.9538	0.9542	0.8920

GENHMM: BRING THE CONCEPT INTO HMM




At time t , the probabilistic model of a state $s \in \mathcal{S}$ is then given by

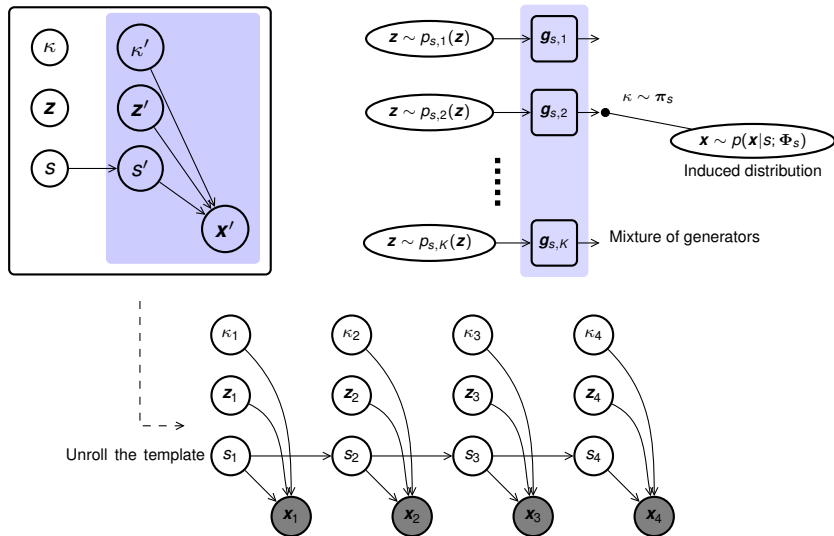
$$p(\mathbf{x}|s; \Phi_s) = \sum_{\kappa=1}^K \pi_{s,\kappa} p(\mathbf{x}|s, \kappa; \theta_{s,\kappa}),$$

where


- $\pi_{s,\kappa} = p(\kappa|s; \mathbf{H})$, naturally $\sum_{\kappa=1}^K \pi_{s,\kappa} = 1$
- $p(\mathbf{x}|s, \kappa; \theta_{s,\kappa})$ is induced by the k th generator $\mathbf{g}_k(\mathbf{z}) = \mathbf{g}(\mathbf{z}; \theta_k)$



GENHMM: BRING THE CONCEPT INTO HMM




ALTERNATIVE VIEW OF GENHMM



Hypothesis set of GenHMM as $\mathcal{H} := \{\mathcal{H} | \mathcal{H} = \{\mathcal{S}, \mathbf{q}, \mathbf{A}, p(\mathbf{x}|s; \Phi_s)\}\}$

- \mathcal{S} : the set of hidden states of \mathcal{H} .
- $\mathbf{q} = [q_1, \dots, q_{|\mathcal{S}|}]^T$: the initial state distributions of \mathcal{H} . $q_i = p(s_1 = i; \mathcal{H})$.
- \mathbf{A} : the transition matrix of states in \mathcal{H} .
- $\Phi = \{\Phi_s | s \in \mathcal{S}\}$.


ALTERNATIVE VIEW OF GENHMM



Hypothesis set of GenHMM as $\mathcal{H} := \{\mathbf{H} | \mathbf{H} = \{\mathcal{S}, \mathbf{q}, \mathbf{A}, p(\mathbf{x}|\mathbf{s}; \Phi_{\mathbf{s}})\}\}$

- \mathcal{S} : the set of hidden states of \mathbf{H} .
- $\mathbf{q} = [q_1, \dots, q_{|\mathcal{S}|}]^T$: the initial state distributions of \mathbf{H} . $q_i = p(s_1 = i; \mathbf{H})$.
- \mathbf{A} : the transition matrix of states in \mathbf{H} .
- $\Phi = \{\Phi_{\mathbf{s}} | \mathbf{s} \in \mathcal{S}\}$.

ALTERNATIVE VIEW OF GENHMM



Hypothesis set of GenHMM as $\mathcal{H} := \{\mathbf{H} | \mathbf{H} = \{\mathcal{S}, \mathbf{q}, \mathbf{A}, p(\mathbf{x}|\mathbf{s}; \Phi_{\mathbf{s}})\}\}$

- \mathcal{S} : the set of hidden states of \mathbf{H} .
- $\mathbf{q} = [q_1, \dots, q_{|\mathcal{S}|}]^T$: the initial state distributions of \mathbf{H} . $q_i = p(s_1 = i; \mathbf{H})$.
- \mathbf{A} : the transition matrix of states in \mathbf{H} .
- $\Phi = \{\Phi_{\mathbf{s}} | \mathbf{s} \in \mathcal{S}\}$.

APPLICATION OF GENHMM

Speech Recognition:

Phoneme classification / recognition

Model	Criterion	K=1	K=3	K=5
GMM-HMM linear variable change	Accuracy	62.3	68.0	68.7
	Precision	67.9	72.6	73.0
	F1	63.7	69.1	69.7
GenHMM non-linear variable change	Accuracy	76.7	77.7	77.7
	Precision	76.9	78.1	78.0
	F1	76.1	77.1	77.0

Robustness to perturbation of noise.

Model	Criterion	White Noise SNR			
		15dB	20dB	25dB	30dB
GMM-HMM	Accuracy	36.6	44.2	50.8	57.1
	Precision	59.2	64.2	68.4	70.6
	F1	39.9	47.7	53.9	59.9
GenHMM	Accuracy	52.4	62.0	69.7	74.3
	Precision	60.0	65.9	71.7	74.8
	F1	52.5	62.0	69.3	73.5

Application to sepsis detection for infants, c.f. see Section 7.5.

- generative training + discriminative training
- innovative feature inspired from acoustic signal feature

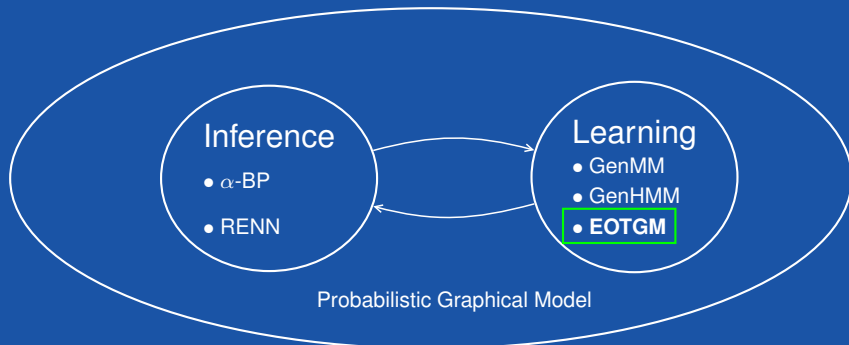
REMARK ON GENMM/GENHMM

Remarks on Learning of GenMM/GenHMM

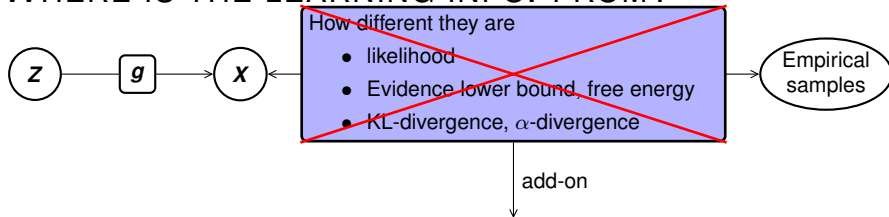
- F , evidence lower bound, equivalent to a negative free energy in DAGs
= \mathcal{Q} in EM + Entropy
- E-step require inference (message-passing for posteriors)
- No optimality in M-step (NN generators, batch-size gradient descent).
- Still, guaranteed non-decreasing $\text{lk}h$. (c.f. Proposition 6.1, Proposition 7.1)

Attributes

- Free dimension for flexibility: number of mixture + complexity of functional form of neural networks
- Tractable likelihood and efficient sampling
- Compatible with classic statistic methods and neural network techniques (error back-propagation, optimizer)



WHERE IS THE LEARNING INFO. FROM?




Optimal transport (OT): moving mass from a dist. to another

$$T(p^*, p) = \min_{\pi \in \Pi(p^*, p)} \left\langle \underbrace{\pi}_{\text{marginalize to } p^*, p}, \underbrace{M}_{\substack{\text{cost matrix} \\ \text{pair-wise sample difference}}} \right\rangle,$$

Key attributes:

- Doesnot require tractible lklh.
- Learning gradient info. from sample comparison
- High complexity, each evaluation is sovling an optimization problem

EOTGM: EOT BASED GENERATIVE MODEL



EOTGM employs:

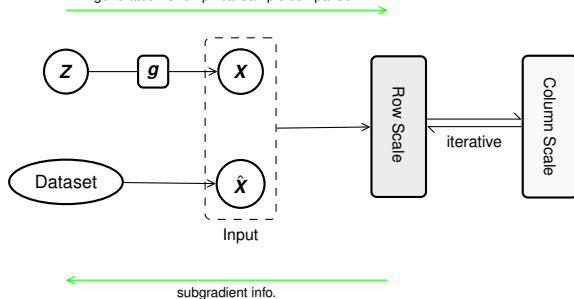
alternatively scaling rows & columns of matrix $e^{-\mathbf{M}/\lambda}$ (Sinkhorn & Knopp)

to extract gradient information (sub-gradient) for generator \mathbf{g}

EOTGM AND EOTGAN

EOTGM

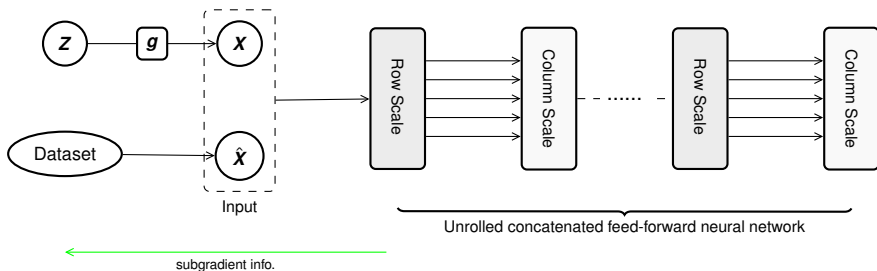
generated v.s. empirical sample comparison



EOTGM AND EOTGAN

EOTGM

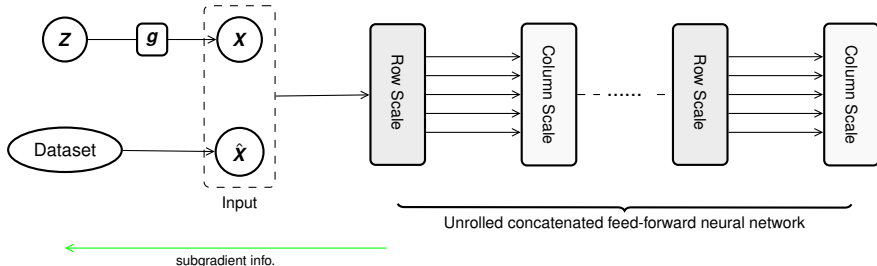
generated v.s. empirical sample comparison



EOTGM AND EOTGAN

EOTGM

generated v.s. empirical sample comparison






EOTGAN: learn implicit distribution with feature mapping, c.f. Section 8.2.2 (Euclidean distance is not suitable for multimedia signals.)

NUMERICAL

→ Toy distribution learning (target at 4-mixture Gaussians) using EOTGM. Real samples (red '+') and contour (red curve), versus generated samples (blue 'o') and contour (blue curve) by g .

↓ Comparison of semantic scores (on MNIST) versus mixing ratio r :

- IS: Inception Score (large is good)
- FID: Frechet Inception Distance (small is good)



SUMMARY

Wrap-up

- Inference with message-passing and analysis
- Inference with free energy minimization by neural networks
- Inference .. Learning: their interactions
- Neural network generators in EM for more flexible modeling; A Further step into temporal models
- A bonus modeling method for likelihood-free learning



Thank you for your attention.
Q&A.